

Placeholder ML Paper (replace with real arXiv PDFs)

Abstract: Retrieval-Augmented Generation (RAG) combines a retriever and a generator to ground language model outputs in external documents. This placeholder paper is used only to test the PDF ingestion step of the pipeline.

1. Introduction

Large language models often hallucinate facts not present in their training data. RAG addresses this by retrieving relevant passages from an external corpus at inference time and conditioning generation on that retrieved context.

2. Method

We use dense embeddings and a vector index to perform nearest neighbor search over document chunks, then pass the top-k chunks to a language model along with the user query.